EEMCS / TCS

AUTOMATING THE CYBERSECURITY TRIAGE PROCESS

A COMPARATIVE STUDY ON THE PERFORMANCE OF LARGE LANGUAGE MODELS

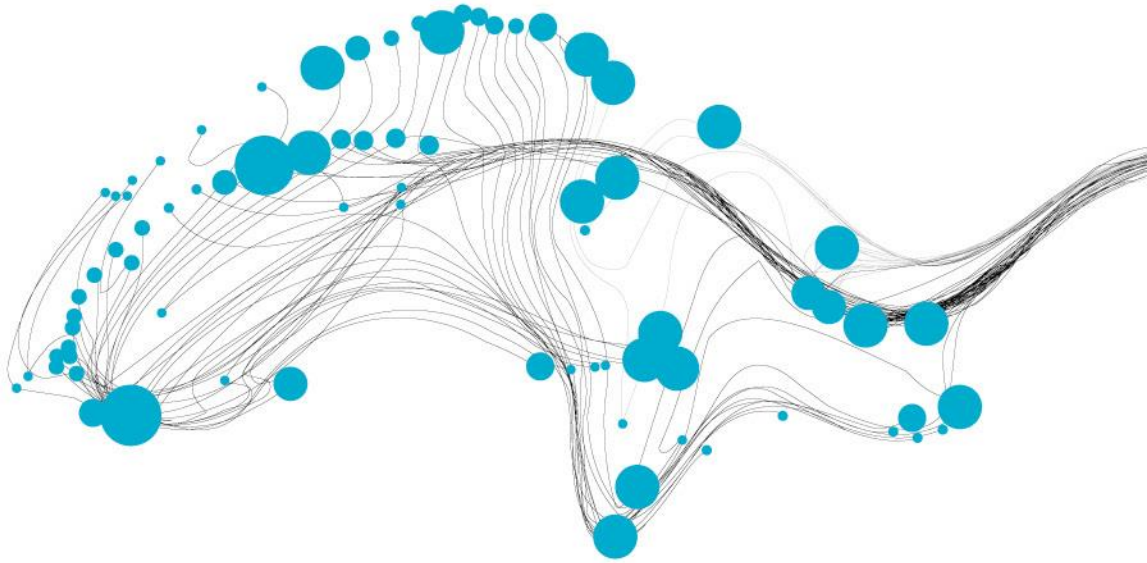

PASCAL BAKKER
SUPERVISED BY JAIR SANTANNA
2024-07-05

UNIVERSITY OF TWENTE.

HOSPITAL – TRIAGE

1. Immediate
2. Emergent
3. Urgent
4. Less urgent
5. Non-urgent

HOSPITAL – TRIAGE

1. Immediate
2. Emergent
3. Urgent
4. Less urgent
5. Non-urgent

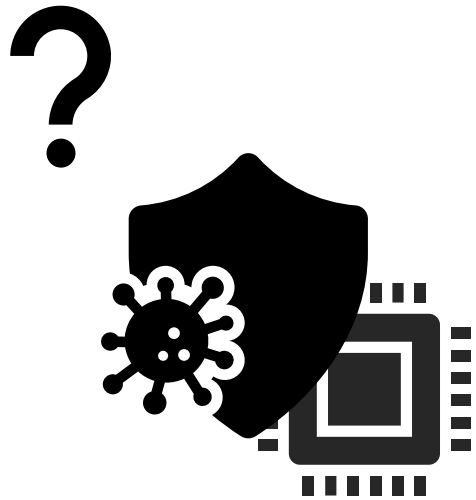

THE PROBLEM WITH TRIAGE

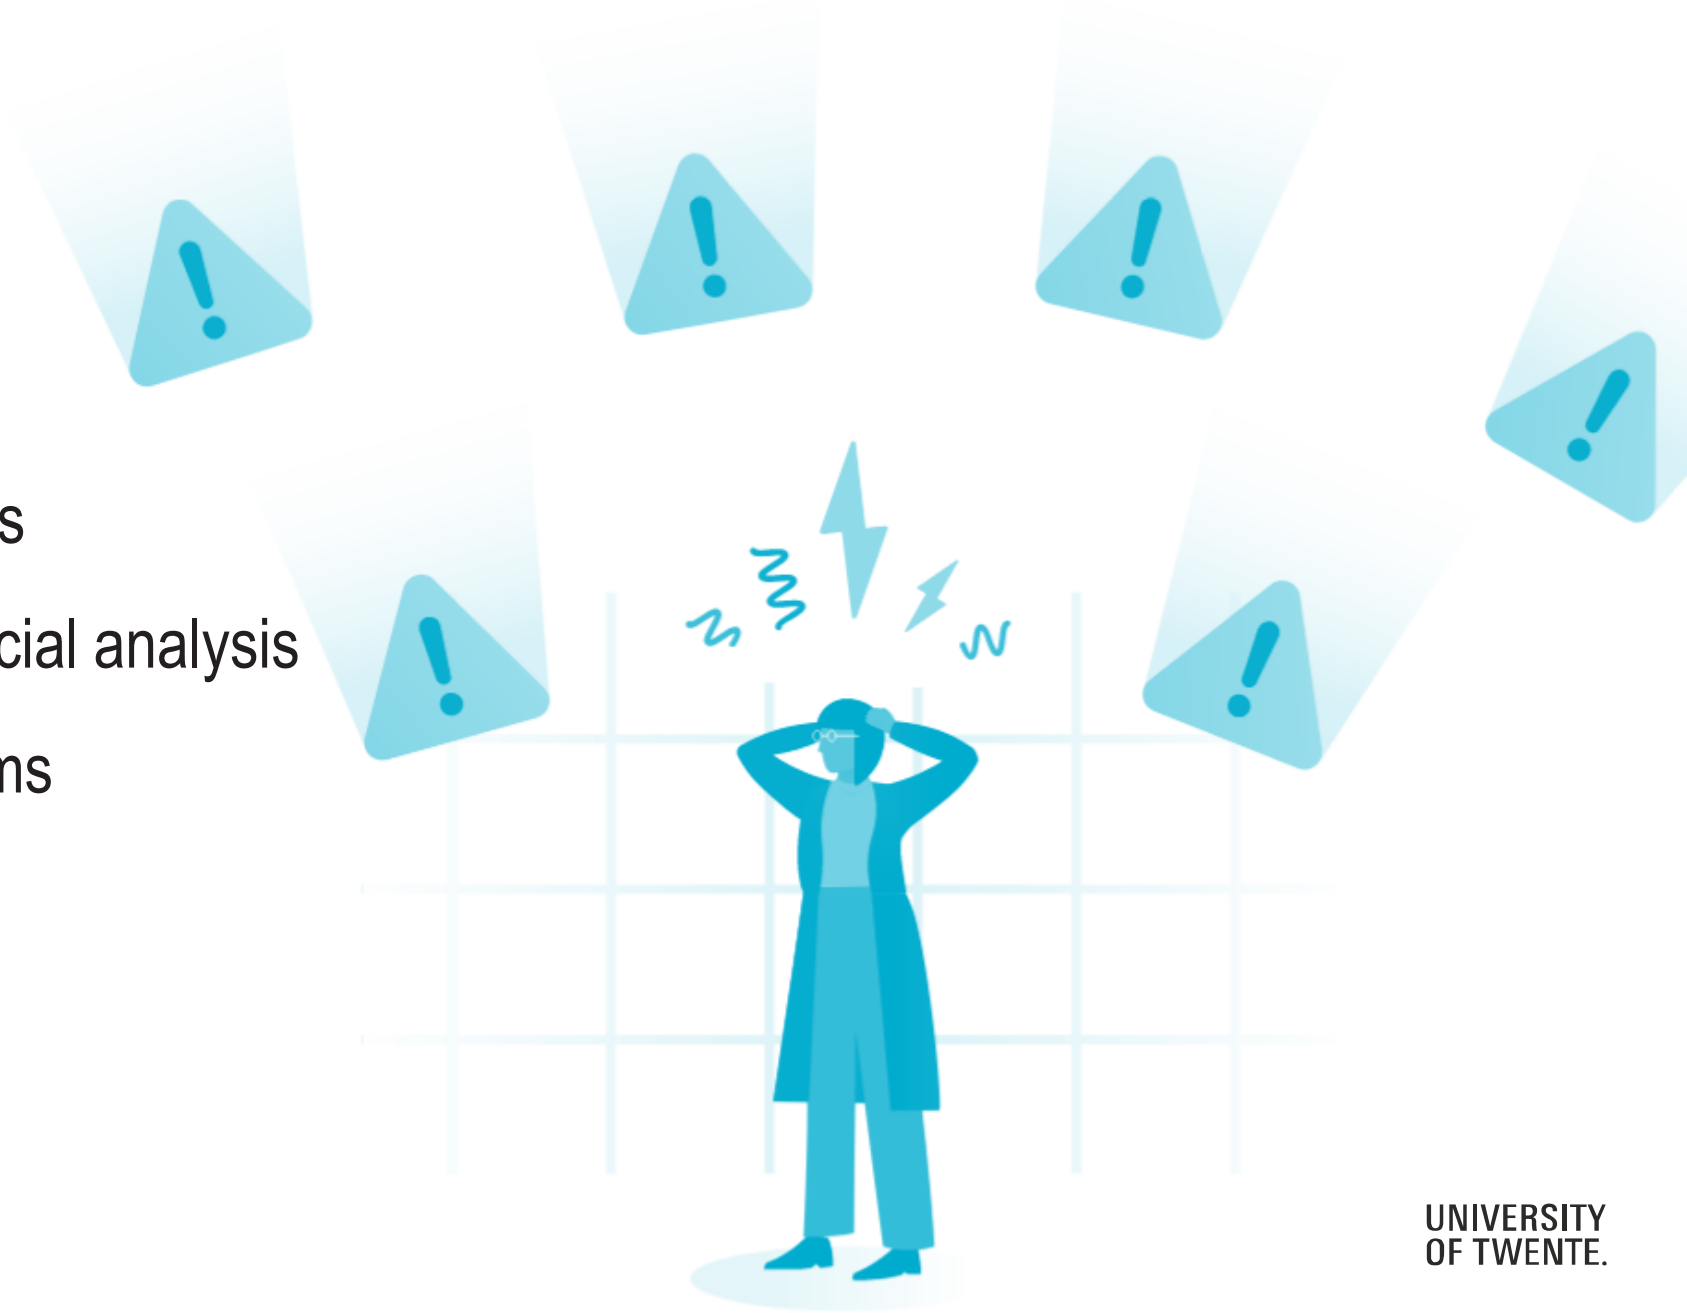

Immense numbers of alarms

- Too little time → Superficial analysis
- Fatigue → Missed alarms
- Burnout / turnover
- Human error

LARGE LANGUAGE MODELS (LLMS)

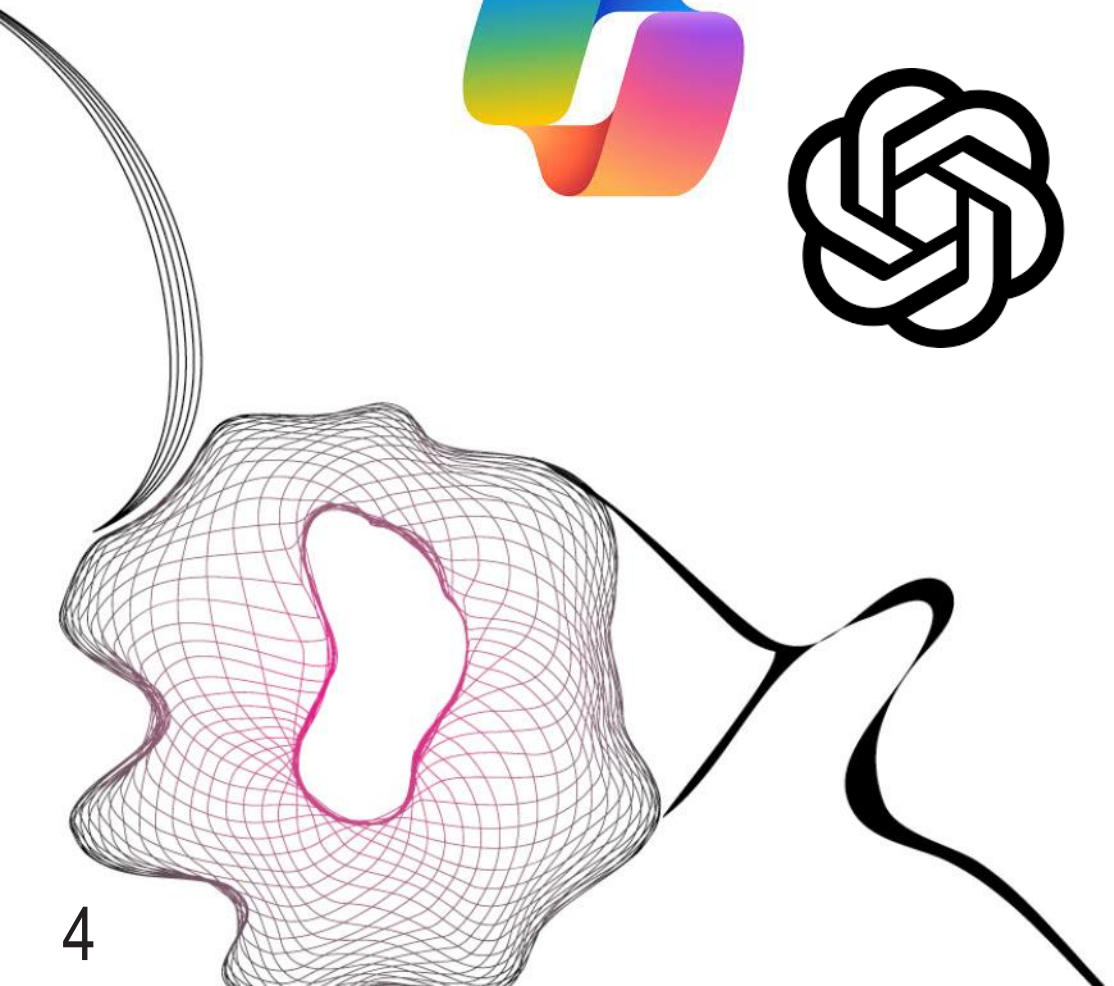

- Natural Language Processing
- Much data → General understanding

1. Generating natural language
2. Identifying contextual relationships
3. Recognizing complex patterns
4. Analyzing semantics

- Many existing applications

→ Automate triage

RESEARCH QUESTIONS

How can LLMs be integrated into the existing incident response workflow to streamline the triage process?

What suitable evaluation metrics should be used to assess the performance of LLMs in cybersecurity triage?

How do different LLMs compare in performance when optimizing the cybersecurity triage process?

DEFINING TRIAGE

Very little documentation

Interview

DEFINING TRIAGE

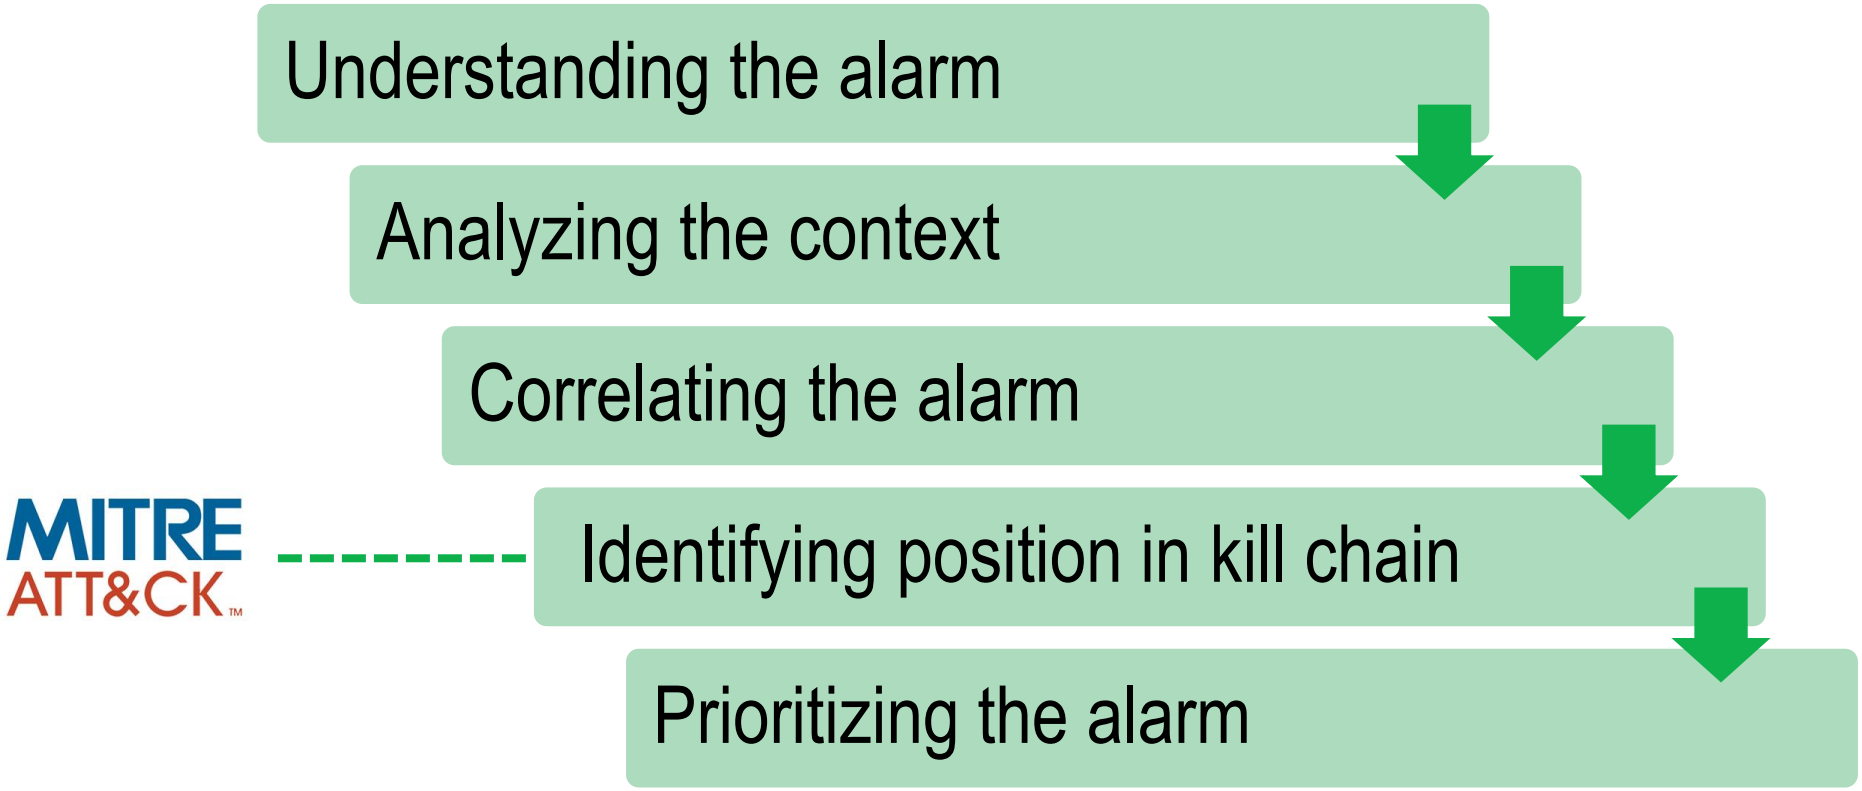

OPTIMIZING TRIAGE
USING LLMS

EXISTING OPTIMIZATIONS

- Organize alarms in trees [12]
- Integrate thread intelligence [41]
- Follow steps of senior analysts [25, 48]

OPTIMIZING TRIAGE
USING LLMS

EXISTING OPTIMIZATIONS

- Organize alarms in trees [12]
- Integrate thread intelligence [41]
- Follow steps of senior analysts [25, 48]

Do not automate natural language tasks

OPTIMIZING TRIAGE
USING LLMS

1. Detecting cybersecurity announcement emails
2. Detecting relation between email and alarm
3. Finding correlation between alarms
4. Determine position in kill chain
5. Determine priority of alarm

OPTIMIZING TRIAGE
USING LLMS

1. Detecting cybersecurity announcement emails
2. Detecting relation between email and alarm
3. Finding correlation between alarms
4. Determine position in kill chain
5. Determine priority of alarm

TASK 1
EVALUATING LLMS

ANNOUNCEMENT DETECTION

1. Give email to LLM
2. Is this a cybersecurity announcement?
3. True / false

TASK 1
EVALUATING LLMS

ANNOUNCEMENT DETECTION

1. Give email to LLM
2. Is this a cybersecurity announcement?
3. True / false

```
Customer X will add the user "sea_line" to the local administrators on the computer "FRLIM-IPC-0017".
```

"True" → ✅

"False" → ❌

TASK 2
EVALUATING LLMS

TACTIC DETECTION

1. Give email to LLM
2. What MITRE ATT&CK tactic can consequential alarms have?
3. e.g. “exfiltration”

TASK 2
EVALUATING LLMS

TACTIC DETECTION

1. Give email to LLM
2. What MITRE ATT&CK tactic can consequential alarms have?
3. e.g. “exfiltration”

```
Customer X will add the user “sea_line” to the local administrators on the computer “FRLIM-IPC-0017”.
```

“privilege escalation” → ✅

“persistence” → ✅

“reconnaissance” → ❌

EVALUATING AND COMPARING LLMS

EVALUATION METRICS

For different prompts:

- F1-score / accuracy
- Median time
- Error rate

EVALUATING AND COMPARING LLMS

EVALUATION METRICS

For different prompts:

- F1-score / accuracy
- Median time
- Error rate

DATASET

- 40 labeled emails
- 20 announcements with labeled tactics

EVALUATING AND COMPARING LLMS

EVALUATION METRICS

For different prompts:

- F1-score / accuracy
- Median time
- Error rate

DATASET

- 40 labeled emails
- 20 announcements with labeled tactics

LARGE LANGUAGE MODELS

(And variations)

- GPT-4
- Llama 3
- Mistral
- Phi-3
- Gemma
- Aya 23
- Code Llama

RESULTS – ANNOUNCEMENT DETECTION

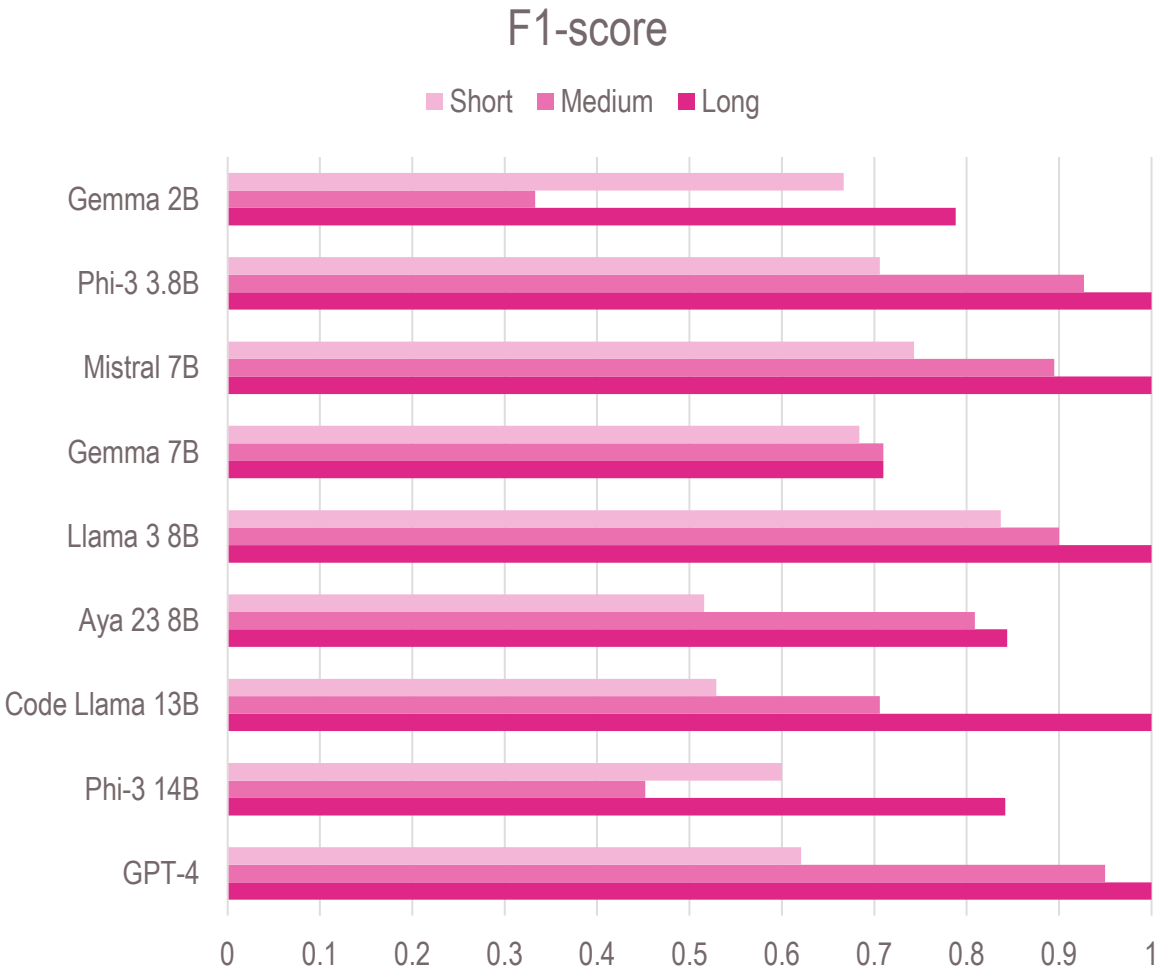

UNIVERSITY OF TWENTE.

RESULTS – ANNOUNCEMENT DETECTION

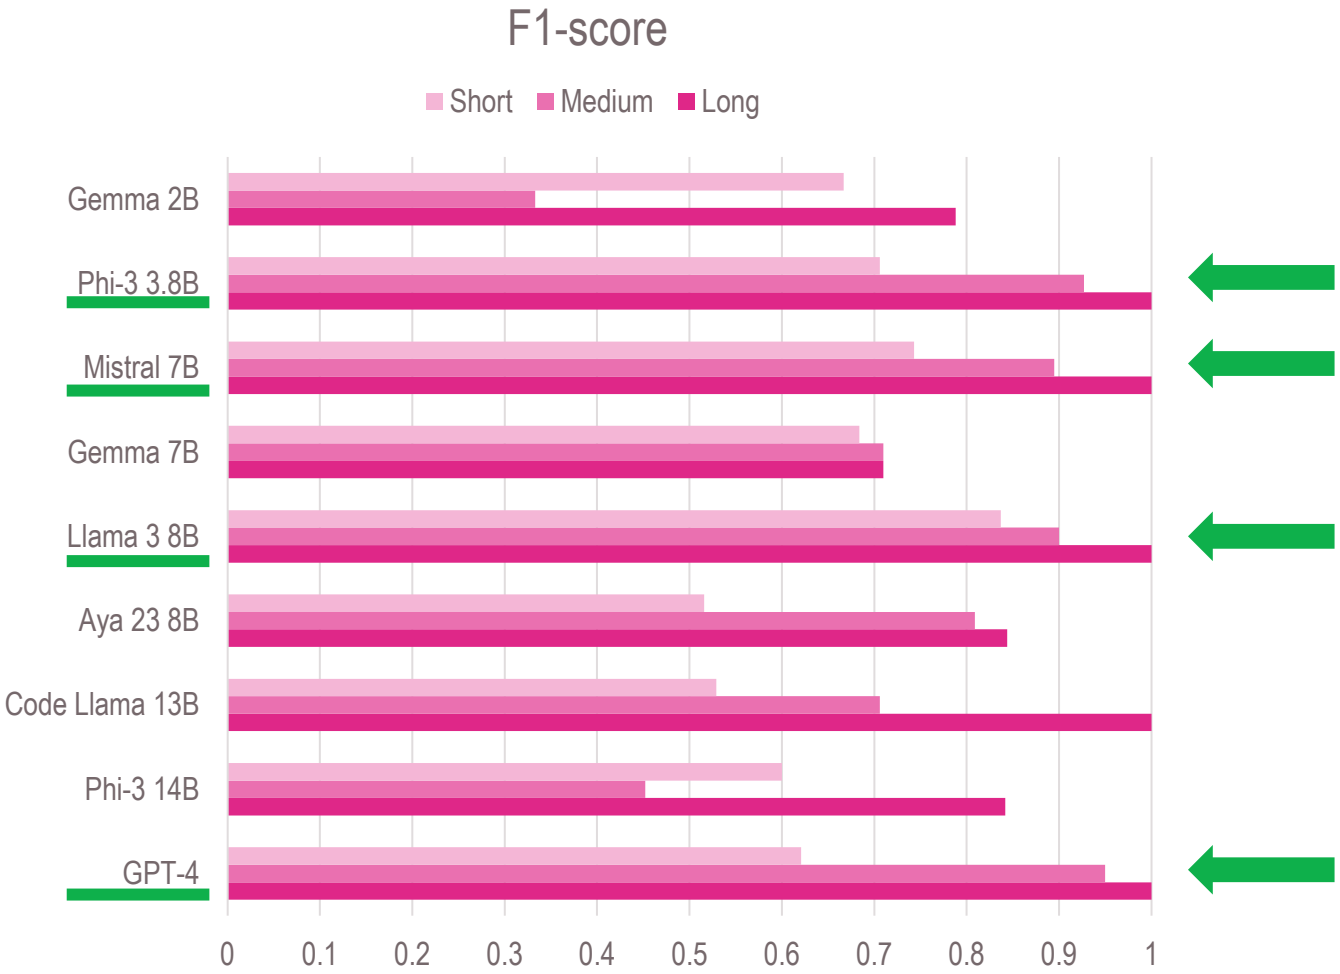

UNIVERSITY OF TWENTE.

RESULTS – ANNOUNCEMENT DETECTION

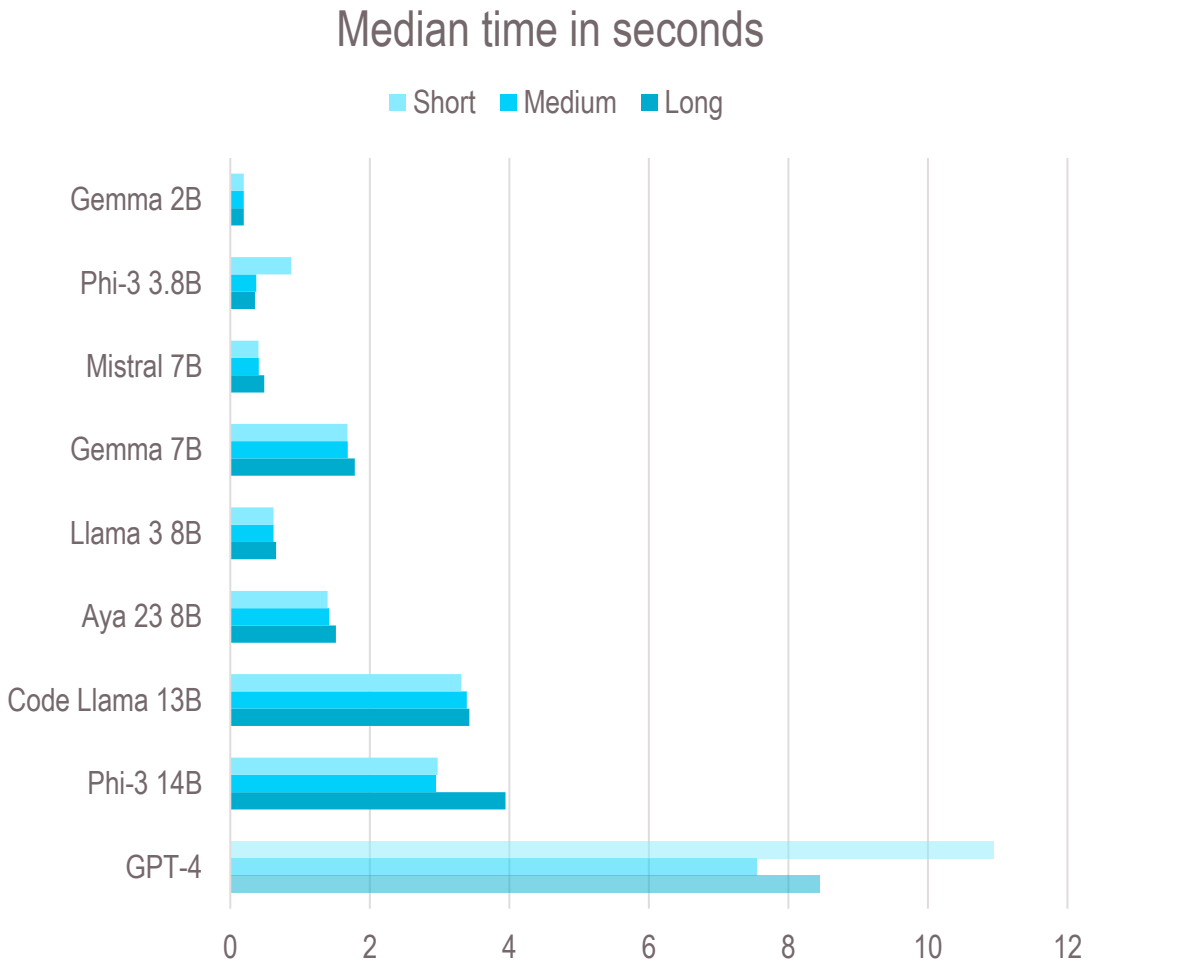

UNIVERSITY OF TWENTE.

RESULTS – ANNOUNCEMENT DETECTION

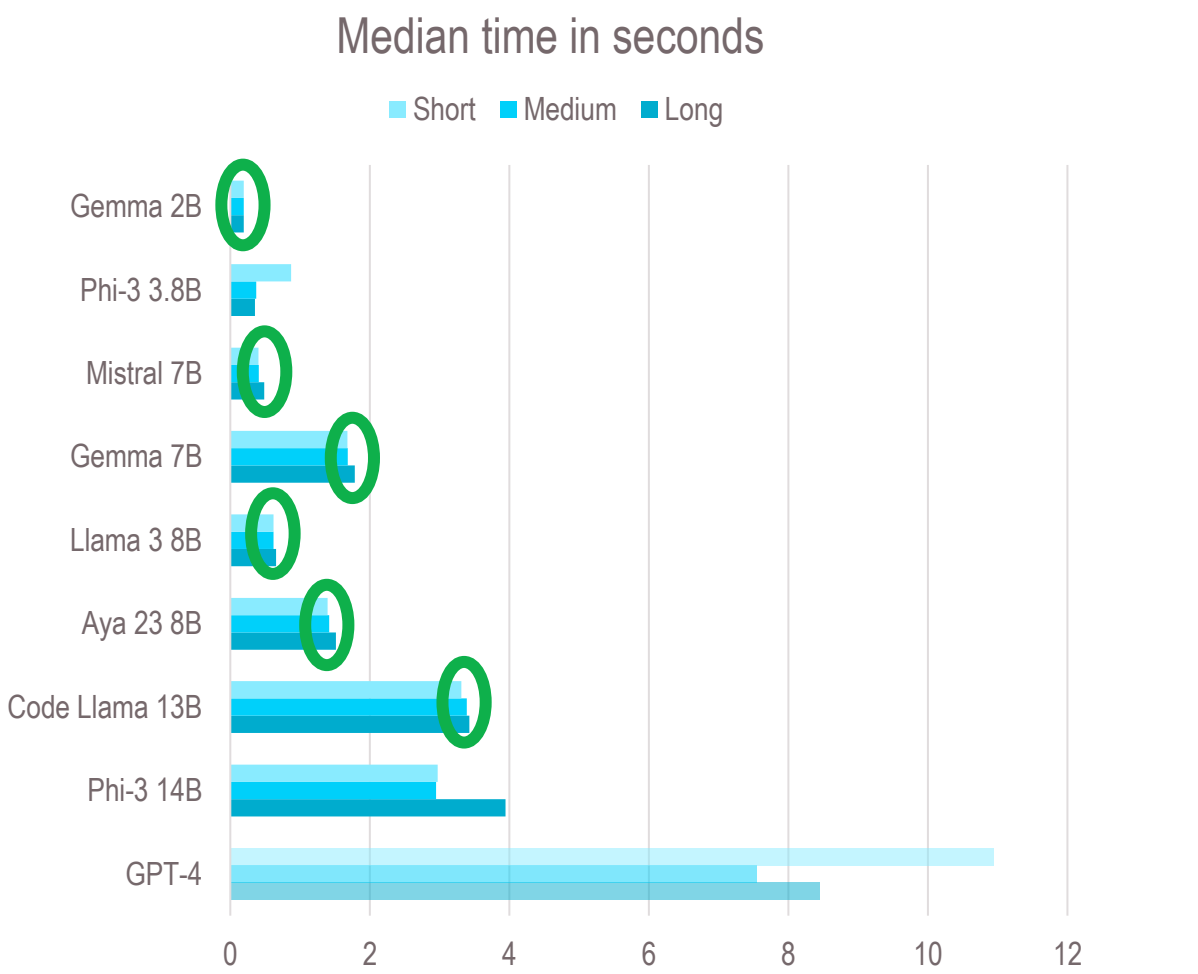

UNIVERSITY OF TWENTE.

RESULTS – ANNOUNCEMENT DETECTION

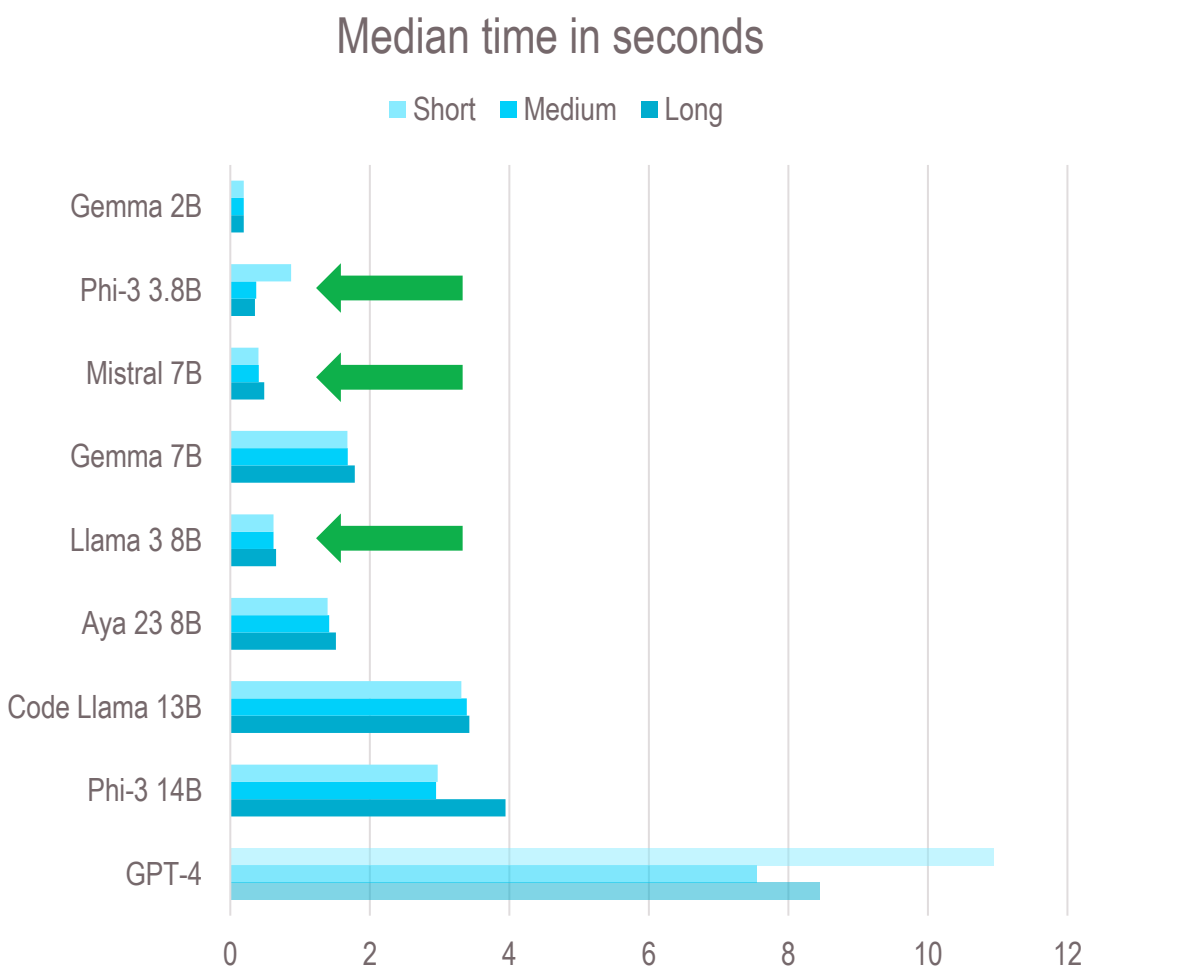

UNIVERSITY OF TWENTE.

RESULTS – TACTIC DETECTION

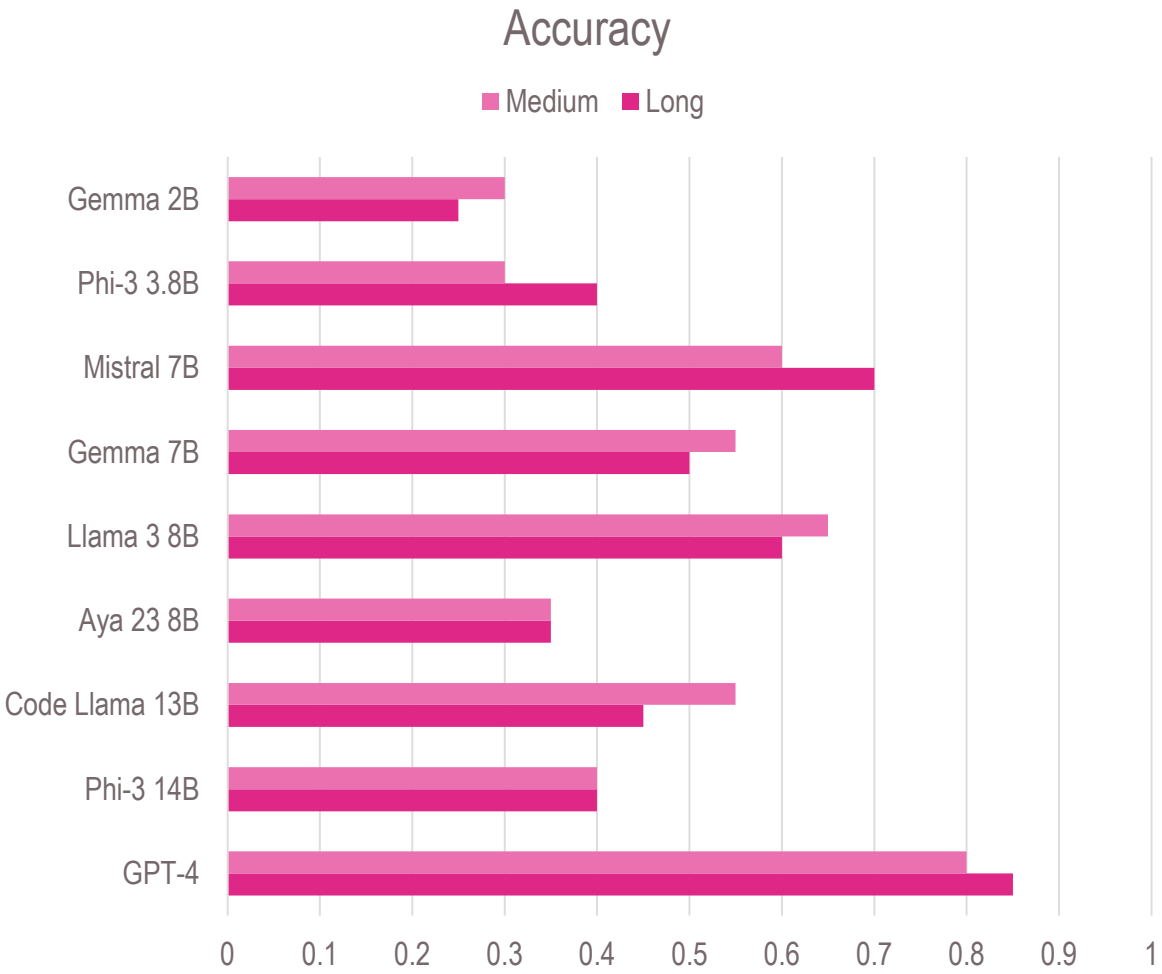

UNIVERSITY OF TWENTE.

RESULTS – TACTIC DETECTION

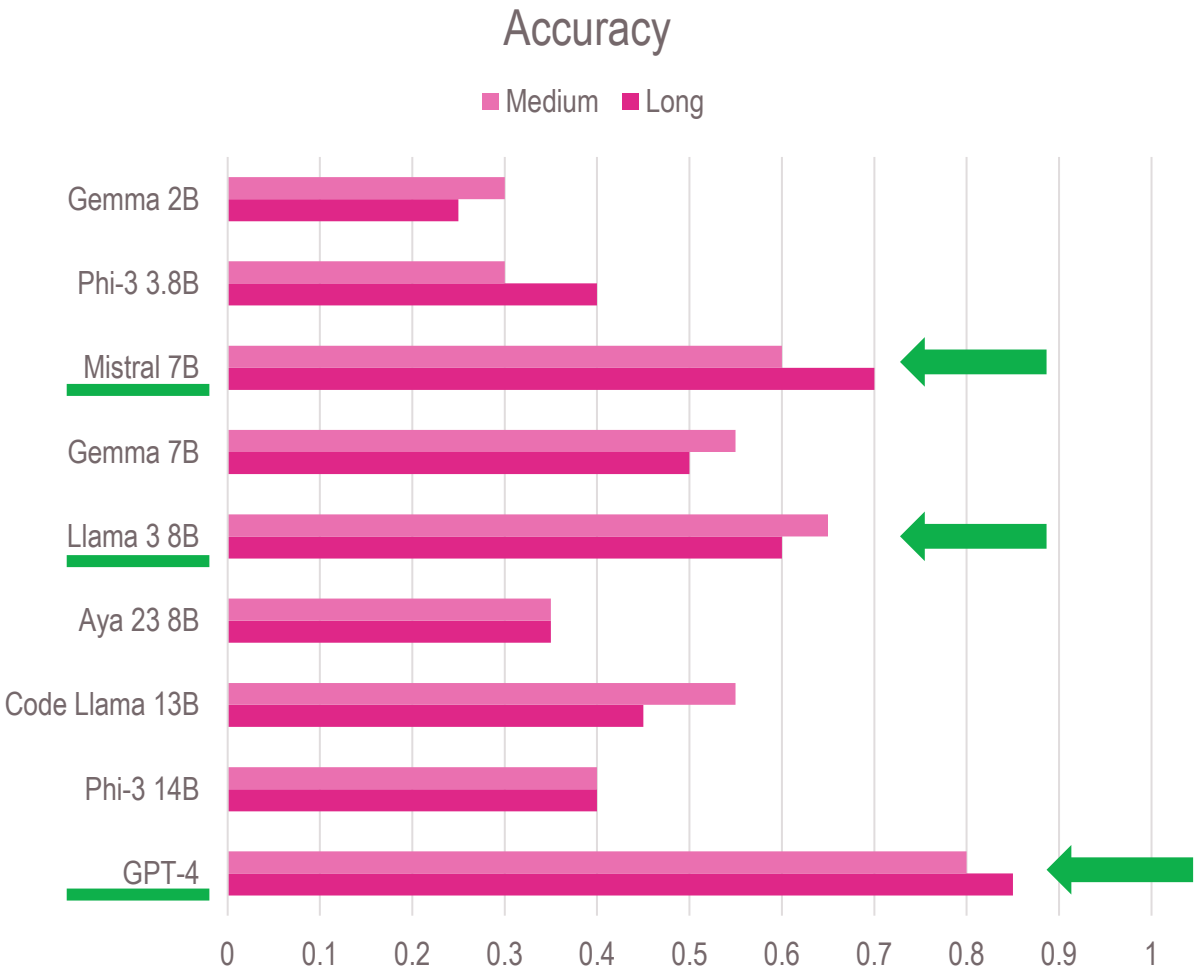

UNIVERSITY OF TWENTE.

RESULTS – TACTIC DETECTION

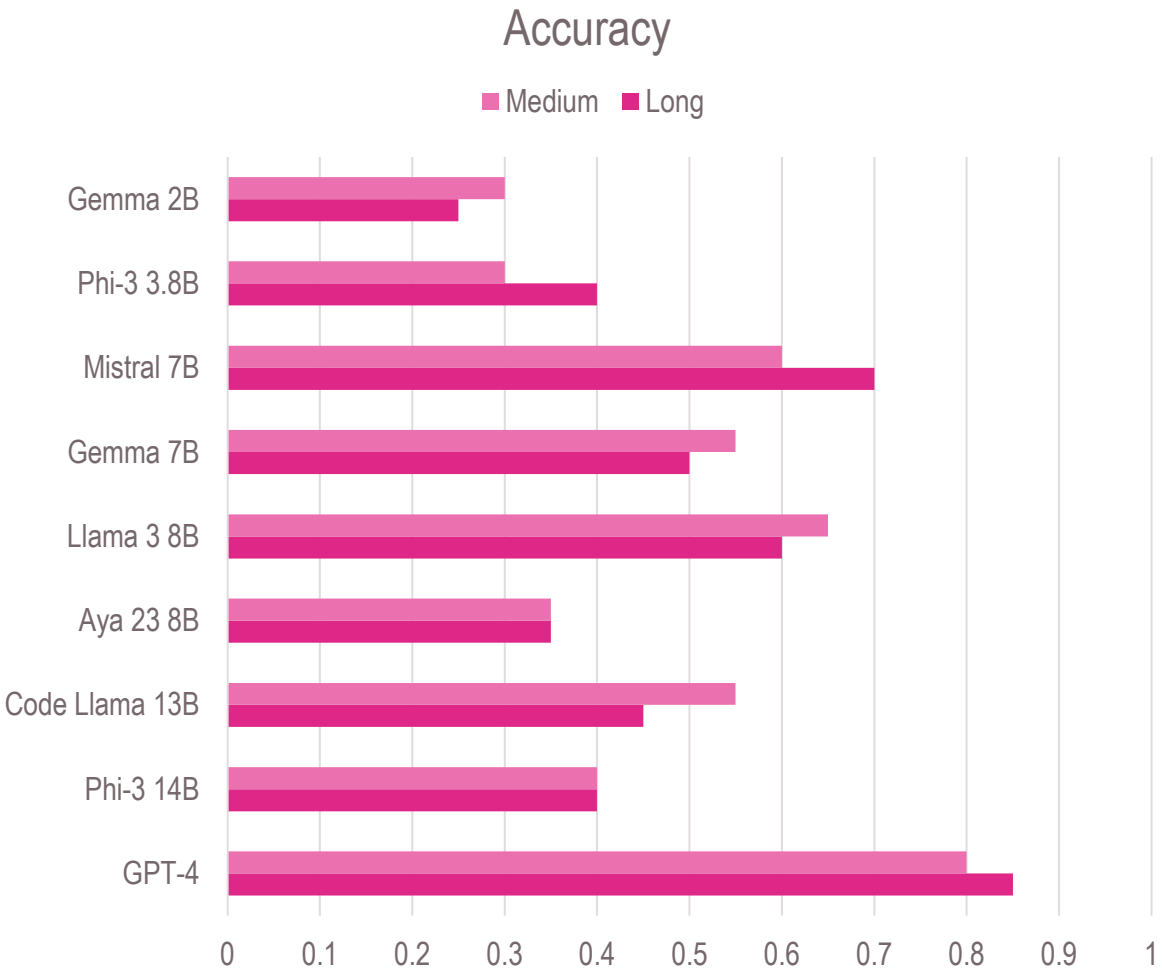

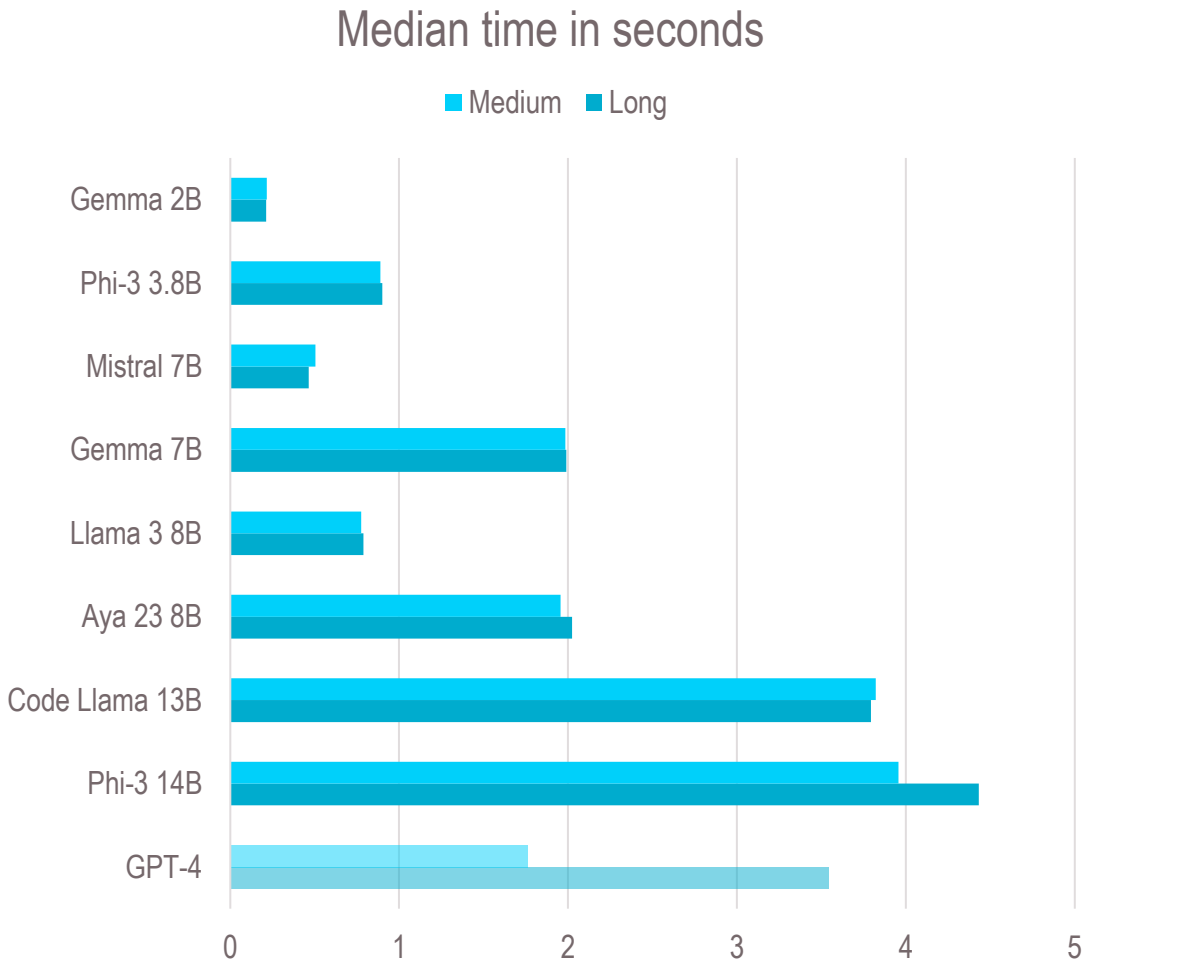

UNIVERSITY OF TWENTE.

RESULTS – TACTIC DETECTION

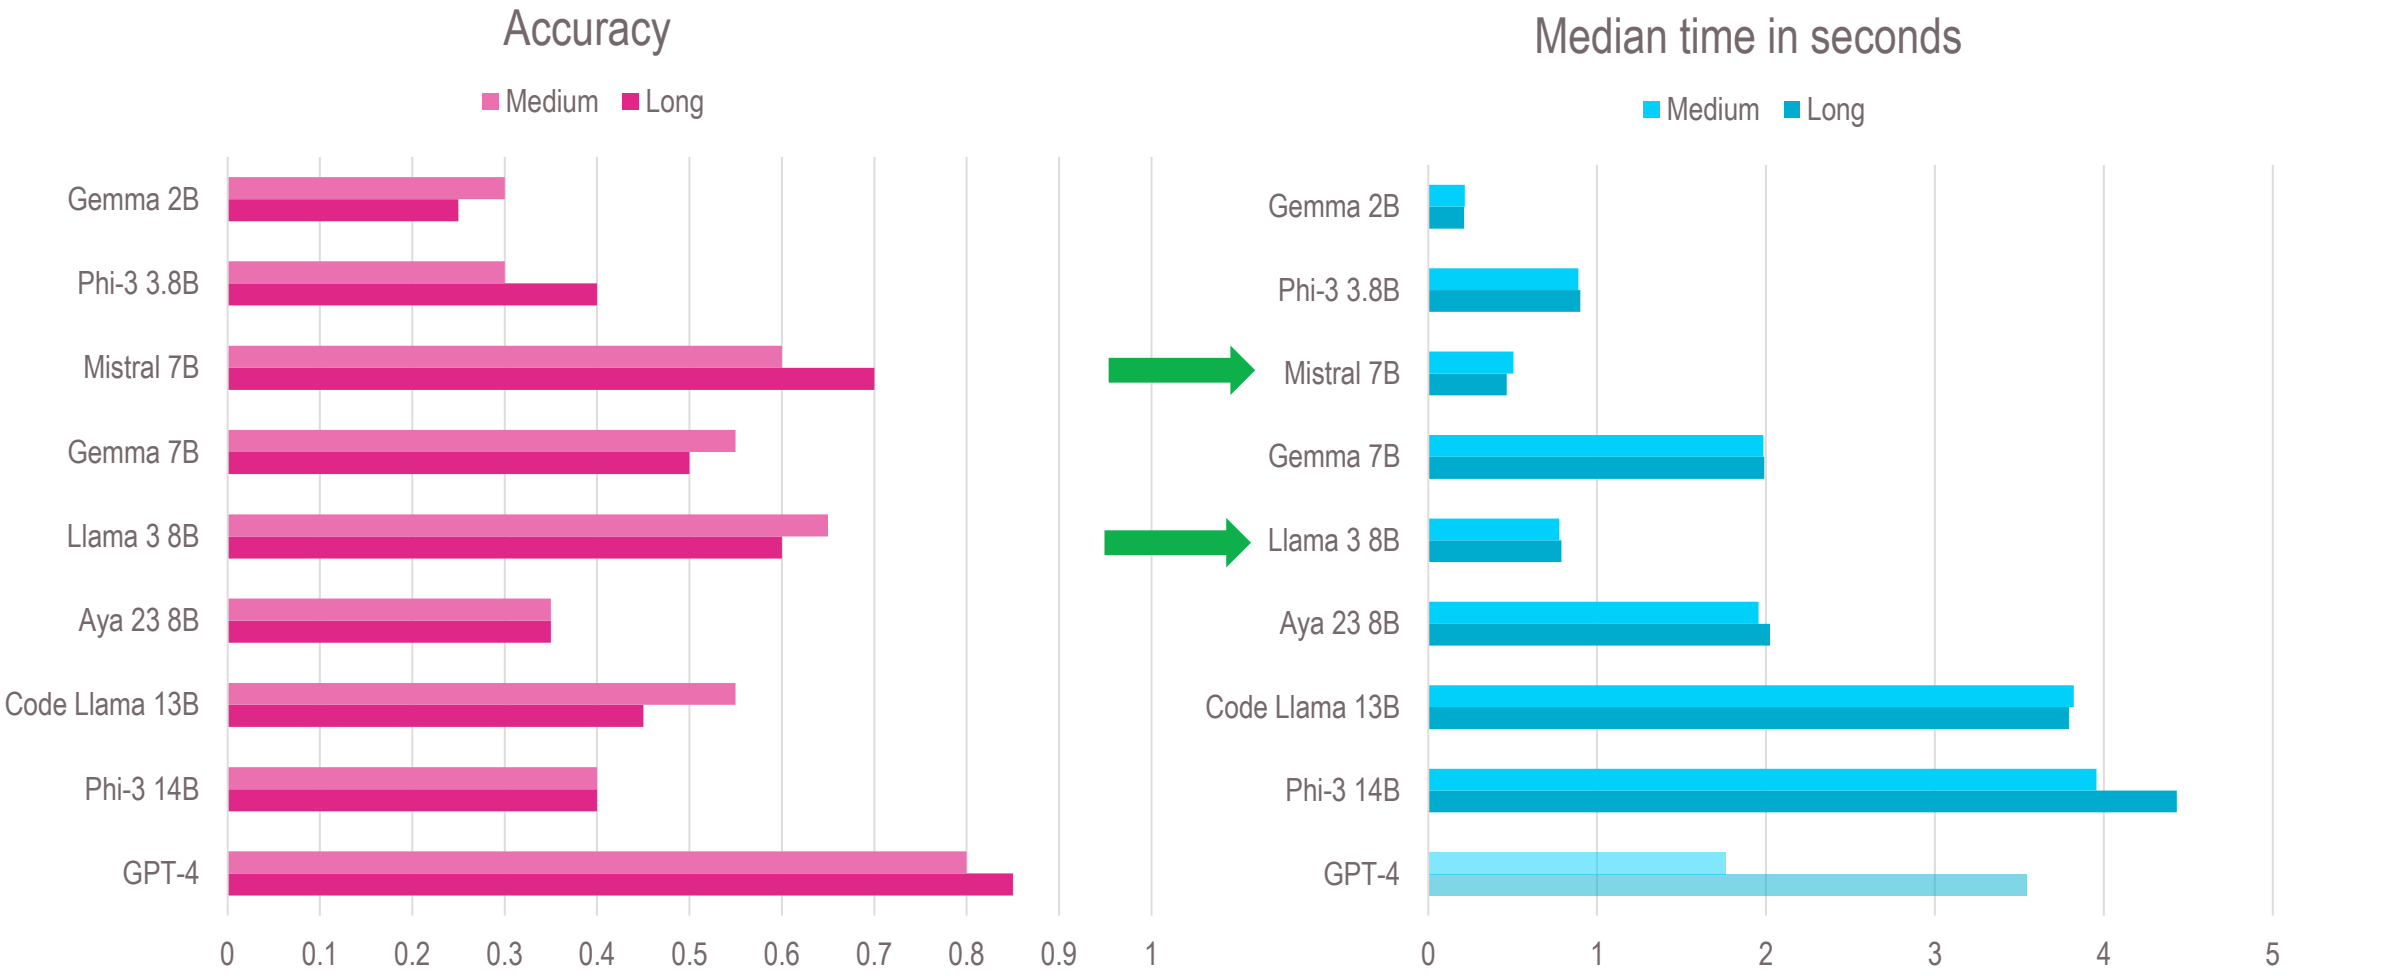

UNIVERSITY OF TWENTE.

RESULTS – COMPARISON

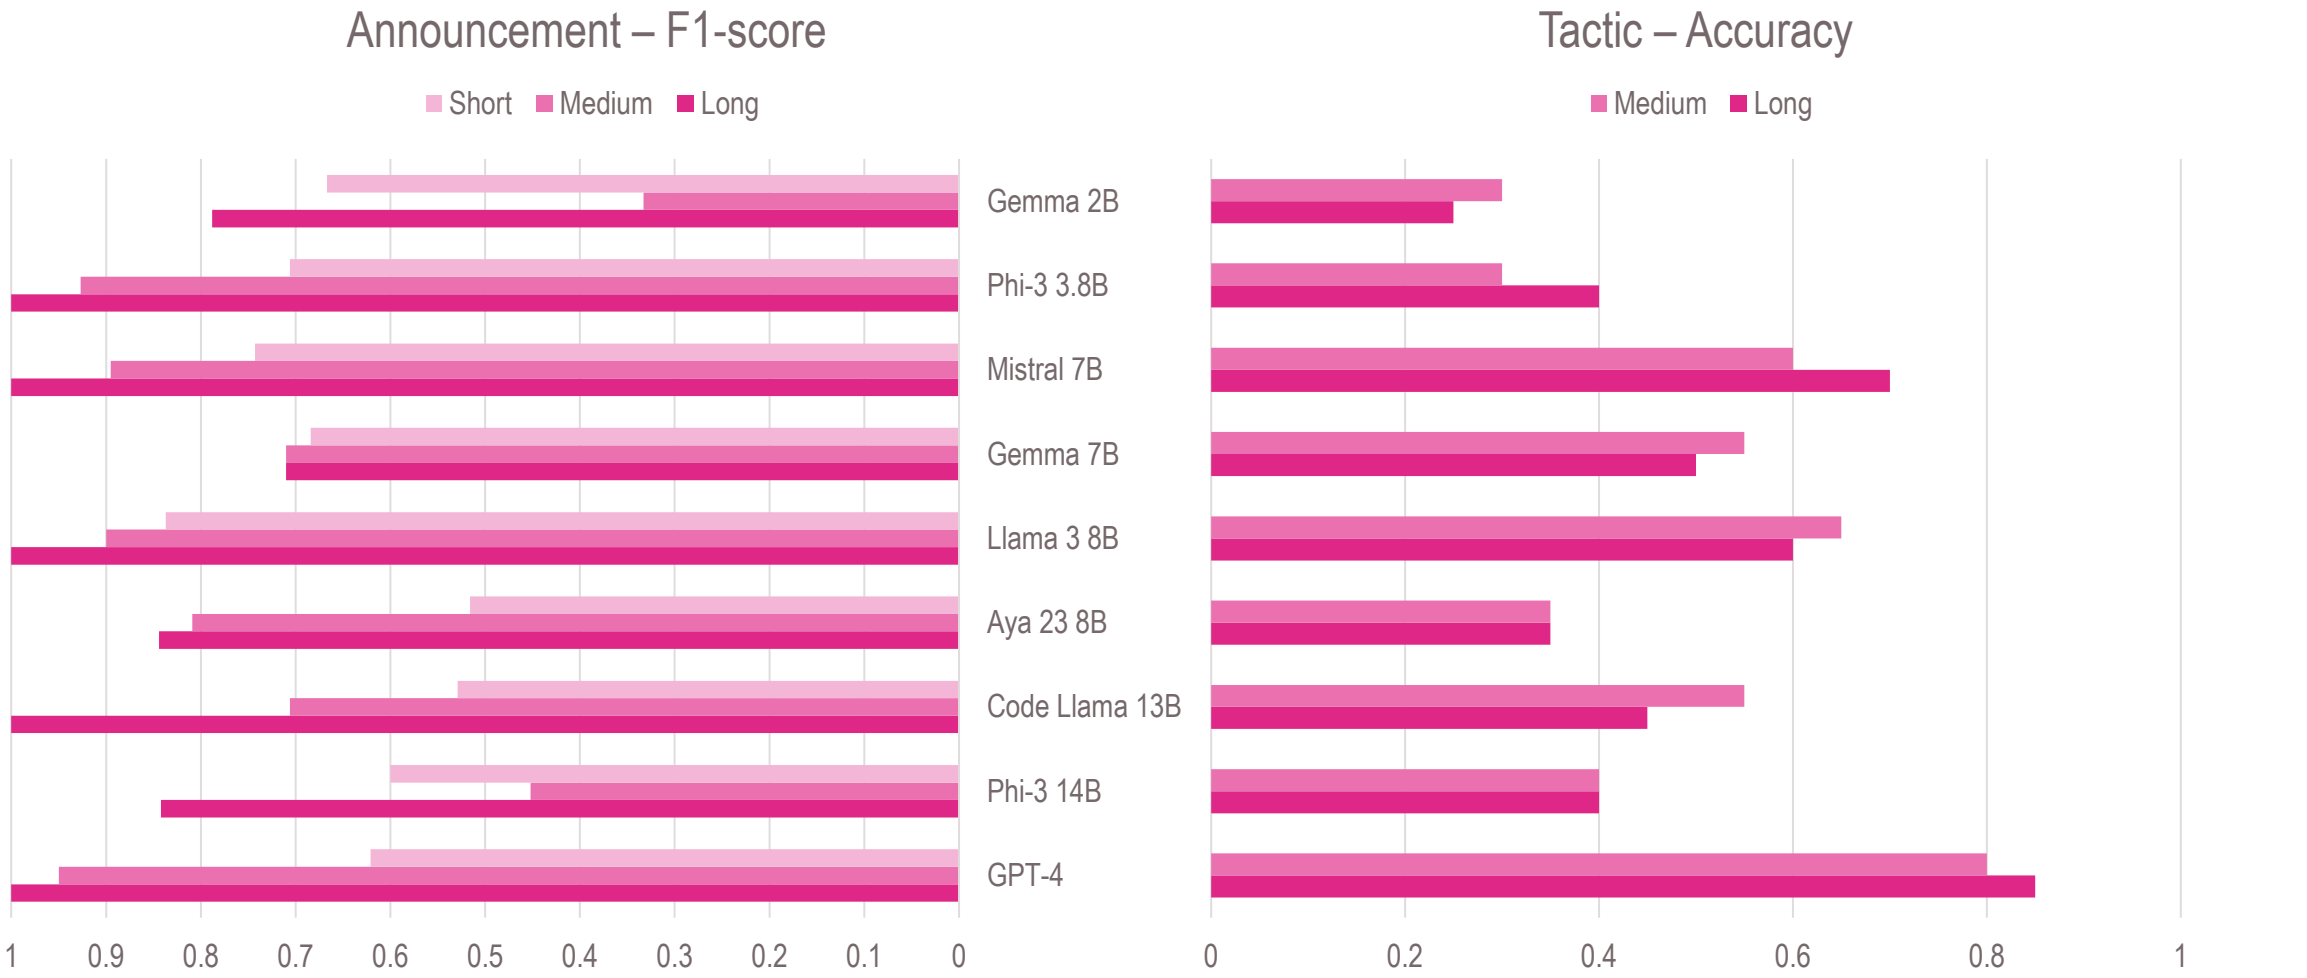

UNIVERSITY OF TWENTE.

RESULTS – COMPARISON

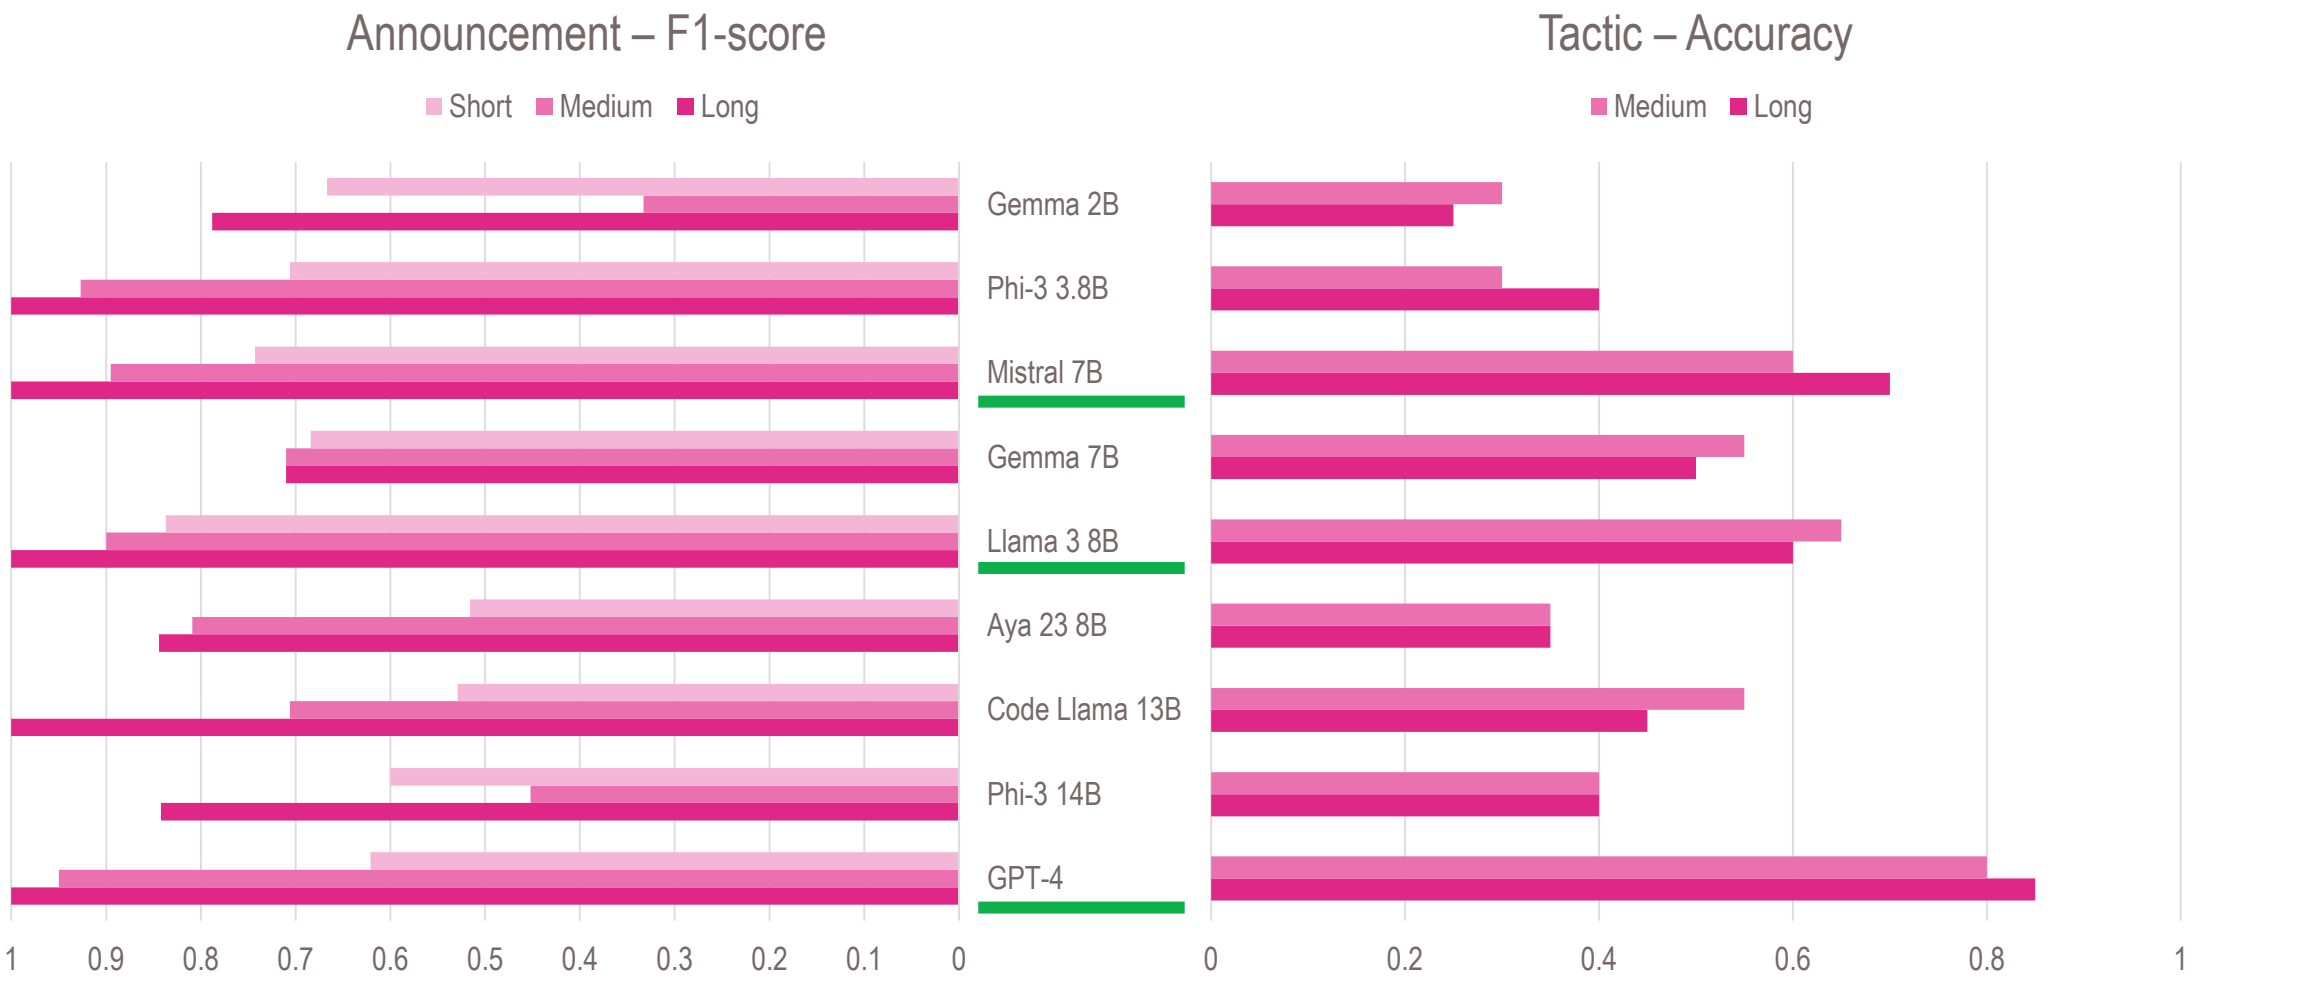

UNIVERSITY OF TWENTE.

RESEARCH QUESTIONS

How can LLMs be integrated into the existing incident response workflow to streamline the triage process?

What suitable evaluation metrics should be used to assess the performance of LLMs in cybersecurity triage?

How do different LLMs compare in performance when optimizing the cybersecurity triage process?

RESEARCH QUESTIONS

How can LLMs be integrated into the existing incident response workflow to streamline the triage process?

What suitable evaluation metrics should be used to assess the performance of LLMs in cybersecurity triage?

How do different LLMs compare in performance when optimizing the cybersecurity triage process?

RESEARCH QUESTIONS

How can LLMs be integrated into the existing incident response workflow to streamline the triage process?

What suitable evaluation metrics should be used to assess the performance of LLMs in cybersecurity triage?

How do different LLMs compare in performance when optimizing the cybersecurity triage process?

RESEARCH QUESTIONS

How can LLMs be integrated into the existing incident response workflow to streamline the triage process?

What suitable evaluation metrics should be used to assess the performance of LLMs in cybersecurity triage?

How do different LLMs compare in performance when optimizing the cybersecurity triage process?

TAKEAWAYS

1. GPT-4 performed the best
   - Followed by Llama 3 and Mistral
   - Phi-3 3.8B good in simple tasks
2. Prompt size had no effect on time
3. Baseline towards further usage of LLMs
   - Defined key steps of triage
   - Identified optimizations
   - Evaluation framework

   → Other tasks and models

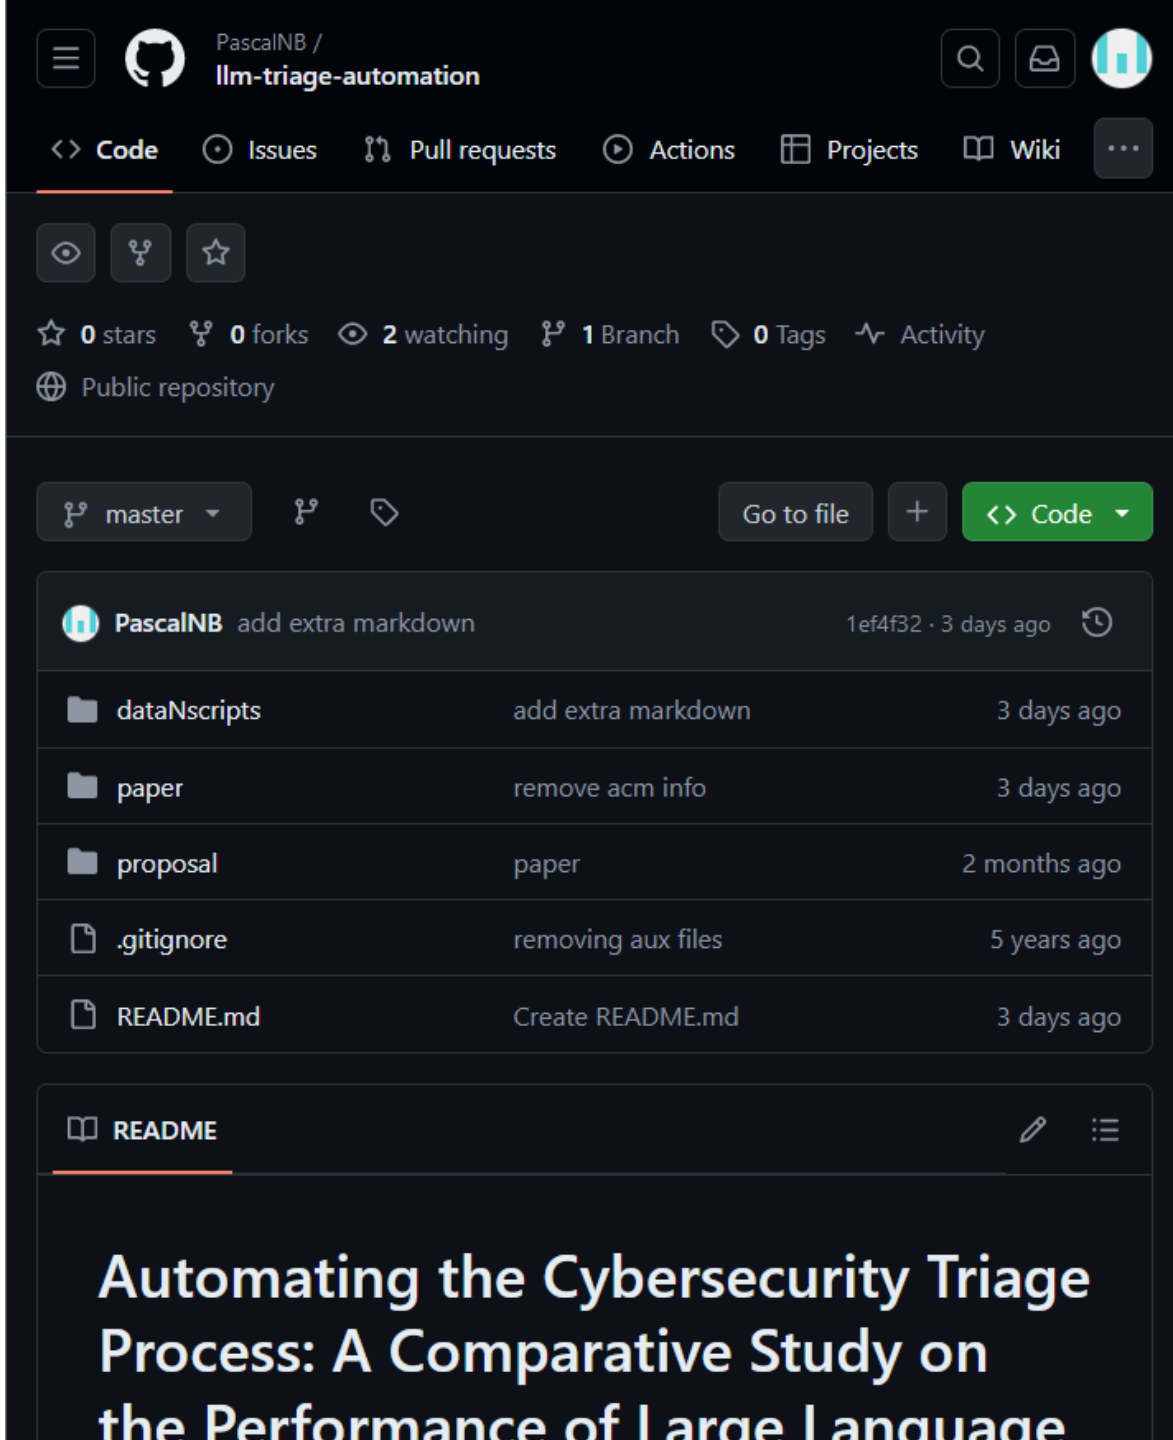

EEMCS / TCS

AUTOMATING THE CYBERSECURITY TRIAGE PROCESS

A COMPARATIVE STUDY ON THE PERFORMANCE OF LARGE LANGUAGE MODELS

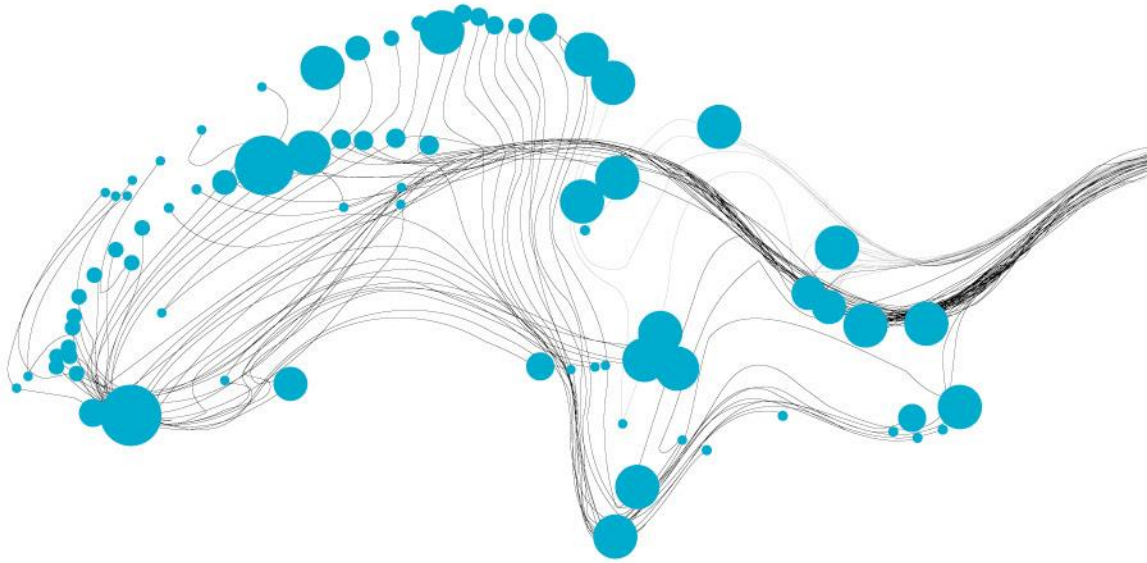

PASCAL BAKKER
SUPERVISED BY JAIR SANTANNA
2024-07-05

UNIVERSITY OF TWENTE.

THE INCIDENT RESPONSE WORKFLOW

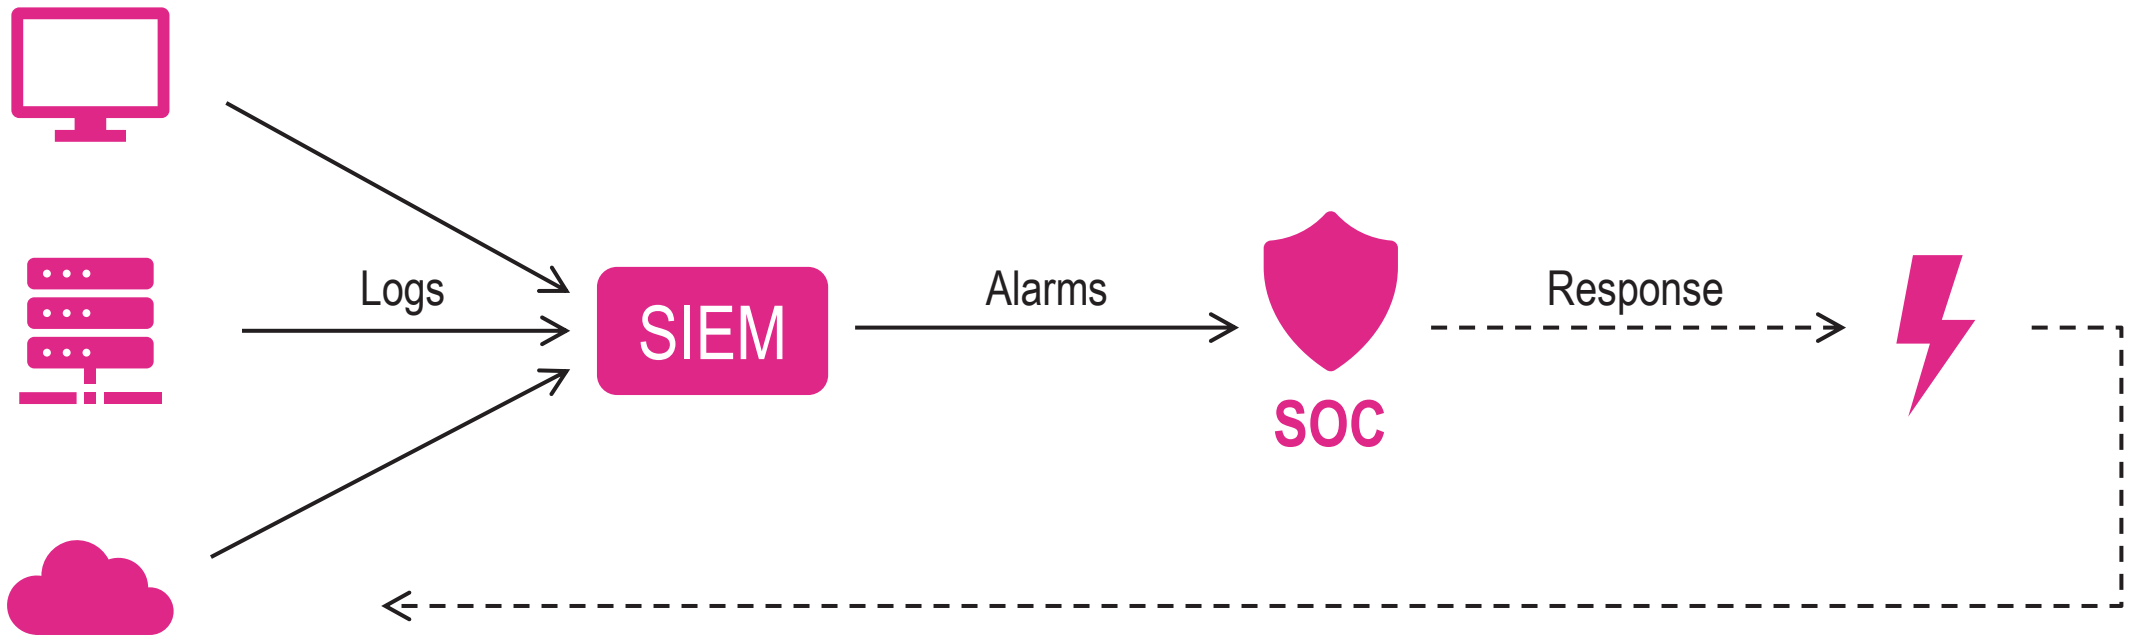

ERROR RATE

```
{
 "is_announcement": True
}
```

```
{
 "is_annoonment": True
}
```

EVALUATION METRICS

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

$$F1 = \frac{2 * Precision * Recall}{Precision + Recall}$$